

Vamsi Krishnam Raju

India, Hyderabad | codewithvamsi401@gmail.com | +91 9390707337

[linkedin.com/in/infinite-coder](https://www.linkedin.com/in/infinite-coder) | www.vamsikrishnamraju.com

SKILLS

Azure Data Factory | Azure Synapse | Azure Databricks | Azure Data Lake | Azure Blob | Apache Spark | PySpark | Spark SQL | Event Hubs | Stream Analytics | Functions | Logic Apps | Azure SQL | Cosmos DB | Data Warehousing | ETL & ELT | Data Modeling | Pipelines | Dimensional Design | DevOps | GitHub | CI/CD Pipelines | Terraform | ARM Infra | Containers (Docker, Kubernetes) | AKS | Infrastructure as Code (IaC) | Monitoring (Azure Monitor, Log Analytics, App Insights) | Power BI | Visualization | SQL | Python | Java | Streaming (Kafka) | Serverless | Security & Access (RBAC) | Data Governance | Lineage & Compliance | Key Vault | Snowflake | Virtual Networking | Performance & Cost Optimization | Automation | Critical Thinking | Team Collaboration | Effective Communication | Flexibility & Time Management | Innovation & Leadership

WORK EXPERIENCE

Data Engineer

August 2022 – October 2024

ReactiveStax(Employer) – Lyft(Client) | Canada

- Engineered a real-time data pipeline using **Databricks Auto Loader**, **PySpark**, and **Delta Live Tables** to continuously ingest, clean, and transform high-volume ride, driver, and payment data from ADLS Gen2, ensuring near real-time analytics.
- Designed scalable, **medallion-architected** ETL pipelines that leverage incremental processing, schema evolution, and custom baseline schemas to maintain consistent data quality across diverse data sources.
- Integrated multiple data sources—including operational ride and driver data from ADLS Gen2 and static lookup/reference data pulled via ADF from GitHub—into a unified **Spark framework** using **DataFrame APIs** and **SparkSQL**.

- Achieved a **50% reduction** in data processing latency by optimizing **Auto Loader** configurations, implementing efficient checkpointing, and tuning Spark clusters for high-throughput processing
- Developed and managed robust data models and **star schema** designs to accurately track ride history, dynamic pricing, and driver performance, facilitating efficient reporting and advanced analytics.
- Streamlined data operations by automating manual processes with **Databricks Workflows**.
- **Enabled real-time decision-making** for actuaries and underwriters by providing insights through a Kibana dashboard, boosting operational efficiency.

Junior Data Engineer

April 2021 – April 2022

ReactiveStax(Employer) – Allstate(Client) | Canada

- Maintained and developed **Azure Databricks** data pipelines for efficient data transformation using Apache Spark.
- **Facilitated seamless data integration** by supporting ingestion and processing tasks across diverse data sources.
- **Developed microservices with Spring Boot**, integrated with Spring Security for robust security, and used Kafka for efficient message brokering.
- **Optimized data storage and retrieval** processes leveraging Azure Data Lake Storage and Blob Storage for high availability.
- **Conducted data analysis** using Spark SQL and DataFrame operations, enabling insightful reporting and analytics.
- **Streamlined data processing** by designing efficient pipelines that reduced transformation times and improved overall performance.
- **Improved data storage mechanisms** to ensure optimal performance, reliability, and scalability in data solutions.

Project Intern

August 2016 – July 2017

Indian Space Research Organisation(ISRO) | India

- **Built a private cloud infrastructure** using OpenStack, configuring essential components to ensure peak performance.
- **Managed network interfaces** and handled name resolution for controller and compute nodes to ensure smooth cloud operations.
- **Developed an intuitive web dashboard** for streamlined management of cloud resources and access controls.
- **Set up and configured KVM Hypervisor**, MariaDB, and RabbitMQ, supporting the stability and scalability of the cloud infrastructure.

- **Laid a strong foundation for future cloud projects**, meeting the organization's research and operational requirements.
- **Enhanced system efficiency and reliability**, contributing to advancements in ISRO's technological capabilities.

PROJECTS

Personal Portfolio

[View in GitHub](#)

- Developed and deployed a personal website using **React**, **Node.js**, JavaScript, HTML, and CSS. Implemented a responsive and interactive user interface, showcasing my portfolio and blog.
- Hosted the website on **AWS Elastic Beanstalk**, ensuring reliable and scalable access. This project demonstrates my full-stack web development and cloud deployment skills.

Microsoft Fabric

[View in GitHub](#)

- Developed a data ingestion pipeline within **Microsoft Fabric** utilising **Azure Data Factory (ADF) v2**.
- Utilized ADF connectors to extract data from diverse sources and orchestrate the landing of data in an ADLS Gen2 lakehouse within the Fabric workspace.

EDUCATION

M.Sc. in Computer Science (Thesis) -
University of Western Ontario (Canada)

August 2018 - Jan 2020

Bachelor of Technology (B. Tech.)
Gitam University, Hyderabad (India)

July 2013 - April 2017